

Data Engineer Technical Assessment

The purpose of this exercise is for you to demonstrate your technical abilities in setting up an end-to-end data pipeline. Some of the technologies we'd like you to use are specific, while others we leave open for interpretation, so you will get to choose.

From a high level, we will provide you with a dataset that you will produce to a message queue, consume that dataset from the message queue and do some manipulations, and send it downstream to a datastore. This should all be done locally via a set of docker containers that communicate with each other.

The deliverables for this assessment are a github repository with the following:

1. A docker compose file that will stand up your data pipeline end-to-end.
2. A set of helper scripts that will help you interact with the data pipeline. Potential helper scripts could do the following:
 - a. Start the data pipeline.
 - b. Stop the data pipeline.
 - c. Produce the events to the data pipeline.
 - d. Monitor the data pipeline.
 - e. Give the status of all of the components of the data pipeline.
3. A README that will document how to use the docker compose file and helper scripts.
4. A design document describing your technical solution and an architecture diagram

Please be ready to explain your pipeline design process, as well as any factors that were taken into consideration.

Instructions:

1. Produce [this](#) dataset to a message queue of your choice using the producer or your choice.
 - a. Potential message queues you could use are RabbitMQ, Redis, Kafka, etc.
 - b. Potential producers you could use are Python, Fluentd, Logstash, etc.
2. Consume the data produced in step 1 from your message queue, manipulate it into the following format, and index the events into an OpenSearch cluster.
 - a. Potential consumers you could use are Python, Fluentd, Logstash, etc.

Unset

```
{
  "time": 1426279439, // epoch time derived from the time field in the event
  "sourcetype": "nginx",
  "index": "nginx",
  "fields": {
```

```
    "region": "us-west-1",
    "assetid": "8972349837489237"
  },
  "event": {
    "remote_ip": "93.180.71.3",
    "remote_user": "-",
    "request": "GET /downloads/product_1 HTTP/1.1",
    "response": 304,
    "bytes": 0,
    "referrer": "-",
    "agent": "Debian APT-HTTP/1.3 (0.8.16~exp12ubuntu10.21)"
  } // this should be all of the data from the event itself, minus time
}
```

3. Using OpenSearch Dashboards, create a dashboard using the data you indexed into OpenSearch. What visualizations you add are up to you, but create something that is insightful for the data set that is provided.